

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
```

```
In [2]: import pandas as pd

df = pd.read_csv("weather.csv")
df.head()
```

Out[2]:

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	WindGustSpeed
0	8.0	24.3	0.0	3.4	6.3	NW	30.0
1	14.0	26.9	3.6	4.4	9.7	ENE	39.0
2	13.7	23.4	3.6	5.8	3.3	NW	85.0
3	13.3	15.5	39.8	7.2	9.1	NW	54.0
4	7.6	16.1	2.8	5.6	10.6	SSE	50.0

5 rows × 22 columns

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 366 entries, 0 to 365
Data columns (total 22 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   MinTemp          366 non-null    float64
 1   MaxTemp          366 non-null    float64
 2   Rainfall         366 non-null    float64
 3   Evaporation     366 non-null    float64
 4   Sunshine         363 non-null    float64
 5   WindGustDir      363 non-null    object  
 6   WindGustSpeed    364 non-null    float64
 7   WindDir9am       335 non-null    object  
 8   WindDir3pm       365 non-null    object  
 9   WindSpeed9am     359 non-null    float64
 10  WindSpeed3pm     366 non-null    int64  
 11  Humidity9am      366 non-null    int64  
 12  Humidity3pm      366 non-null    int64  
 13  Pressure9am      366 non-null    float64
 14  Pressure3pm      366 non-null    float64
 15  Cloud9am         366 non-null    int64  
 16  Cloud3pm         366 non-null    int64  
 17  Temp9am          366 non-null    float64
 18  Temp3pm          366 non-null    float64
 19  RainToday         366 non-null    object  
 20  RISK_MM          366 non-null    float64
 21  RainTomorrow     366 non-null    object  
dtypes: float64(12), int64(5), object(5)
memory usage: 63.0+ KB
```

```
In [4]: df.describe()
```

	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustSpeed	W
count	366.000000	366.000000	366.000000	366.000000	363.000000	364.000000	
mean	7.265574	20.550273	1.428415	4.521858	7.909366	39.840659	
std	6.025800	6.690516	4.225800	2.669383	3.481517	13.059807	
min	-5.300000	7.600000	0.000000	0.200000	0.000000	13.000000	
25%	2.300000	15.025000	0.000000	2.200000	5.950000	31.000000	
50%	7.450000	19.650000	0.000000	4.200000	8.600000	39.000000	
75%	12.500000	25.500000	0.200000	6.400000	10.500000	46.000000	
max	20.900000	35.800000	39.800000	13.800000	13.600000	98.000000	



```
In [5]: df.shape
```


```
Out[5]: (366, 22)
```

```
In [6]: df.columns
```


```
Out[6]: Index(['MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation', 'Sunshine',
       'WindGustDir', 'WindGustSpeed', 'WindDir9am', 'WindDir3pm',
       'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am', 'Humidity3pm',
       'Pressure9am', 'Pressure3pm', 'Cloud9am', 'Cloud3pm', 'Temp9am',
       'Temp3pm', 'RainToday', 'RISK_MM', 'RainTomorrow'],
      dtype='object')
```

```
In [7]: import seaborn as sns
import matplotlib.pyplot as plt


sns.pairplot(df)
plt.show()
```




```
In [8]: plt.figure(figsize=(8,5))
plt.scatter(df['MaxTemp'], df['Humidity3pm'])
plt.xlabel('Max Temperature')
plt.ylabel('Humidity at 3pm')
plt.title('Max Temperature vs Humidity (3pm)')
plt.show()
```



```
In [9]: sns.pairplot(df[['MinTemp', 'MaxTemp', 'Rainfall']])
plt.show()
```




```
In [10]: plt.figure(figsize=(8,5))
plt.scatter(df['MinTemp'], df['Humidity9am'])
plt.xlabel('Min Temperature')
plt.ylabel('Humidity at 9am')
plt.title('Min Temperature vs Humidity (9am)')
plt.show()
```



```
In [11]: numeric_df = df.select_dtypes(include='number')

plt.figure(figsize=(12,8))
sns.heatmap(numeric_df.corr(), annot=False, cmap="coolwarm")
plt.title("Correlation Heatmap of Weather Data")
plt.show()
```



```
In [12]: # Select only numeric columns
numeric_df = df.select_dtypes(include='number')

# Separate features and target
X = numeric_df.drop('Rainfall', axis=1)
y = numeric_df['Rainfall']

X.shape, y.shape
```

```
Out[12]: ((366, 16), (366,))
```

```
In [13]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

X_train.shape, X_test.shape
```

```
Out[13]: ((292, 16), (74, 16))
```

```
In [15]: numeric_df.isnull().sum()
```

```
Out[15]: MinTemp      0
          MaxTemp      0
          Rainfall      0
          Evaporation   0
          Sunshine      3
          WindGustSpeed  2
          WindSpeed9am   7
          WindSpeed3pm   0
          Humidity9am    0
          Humidity3pm    0
          Pressure9am    0
          Pressure3pm    0
          Cloud9am       0
          Cloud3pm       0
          Temp9am        0
          Temp3pm        0
          RISK_MM         0
          dtype: int64
```

```
In [16]: import numpy as np
np.isinf(numeric_df).sum()
```

```
Out[16]: MinTemp      0
          MaxTemp      0
          Rainfall      0
          Evaporation   0
          Sunshine      0
          WindGustSpeed 0
          WindSpeed9am   0
          WindSpeed3pm   0
          Humidity9am    0
          Humidity3pm    0
          Pressure9am    0
          Pressure3pm    0
          Cloud9am       0
          Cloud3pm       0
          Temp9am        0
          Temp3pm        0
          RISK_MM        0
          dtype: int64
```

```
In [17]: from sklearn.impute import SimpleImputer

imputer = SimpleImputer(strategy='mean')

X_train_imputed = imputer.fit_transform(X_train)
X_test_imputed = imputer.transform(X_test)
```

```
In [18]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error

model = LinearRegression()
model.fit(X_train_imputed, y_train)

y_pred = model.predict(X_test_imputed)

r2 = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)

r2, mse
```

```
Out[18]: (0.19041298786157834, 34.2256782862109)
```

```
In [21]: X = df[['MinTemp', 'MaxTemp']]
y = df['Rainfall']
```

```
In [22]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_
```

```
In [24]: model = LinearRegression()
model.fit(X_train, y_train)
```

```
Out[24]: ▾ LinearRegression ⓘ ?
```

► Parameters

```
In [25]: y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
print(f'Mean Squared Error for Rainfall Prediction: {mse}')
```

Mean Squared Error for Rainfall Prediction: 37.0768456005826

```
In [27]: df.groupby("RainToday")["MaxTemp"].mean()
```

```
Out[27]: RainToday
No      20.756667
Yes     19.612121
Name: MaxTemp, dtype: float64
```

```
In [28]: df.groupby("RainTomorrow")["MaxTemp"].mean()
```

```
Out[28]: RainTomorrow
No      20.396000
Yes     21.251515
Name: MaxTemp, dtype: float64
```